

CS35L – Fall 2018

Slide set:	4.2
Slide topics:	GDB, debugging
Assignment:	4

Debugging

- Finding and eliminating errors from programs
- Grace Hopper and the “First actual case of bug being found”

92
9/9

0800 Antran started { 1.2700 9.037 847 025
1000 " stopped - antran ✓ 9.037 846 995 correct
13° UC (033) MP - MC 1.130476415 (2) 4.615925059 (-2)
(033) PRO 2 2.130476415
correct 2.130476415
Relays 6-2 in 033 failed special speed test
in Relay " 11.00 test.
Relay #70 Panel F
(moth) in relay.

1100 Started Cosine Tape (Sine check)
1525 Started Multi Adder Test.

1545 First actual case of bug being found.
163100 Antran started.
1700 closed down.

Relay 2145
Relay 3376



Debugging Process

- Reproduce the bug
- Simplify program input
- Use a debugger to track down the origin of the problem
- Fix the problem

Debugger

- A program that is used to run and debug other (target) programs
- Advantages:

Programmer can:

- step through source code line by line
 - each line is executed on demand
- interact with and inspect program at run-time
- If program crashes, the debugger outputs where and why it crashed

GDB – GNU Debugger

- Debugger for several languages
 - C, C++, Java, Objective-C... more
- Allows you to inspect what the program is doing at a certain point during execution
- Logical errors and segmentation faults are easier to find with the help of gdb

Using GDB

1. Compile Program

- Normally: \$ gcc [flags] <source files> -o <output file>
- Debugging: \$ gcc [other flags] **-g** <source files> -o <output file>
 - enables built-in debugging support

2. Specify Program to Debug

- \$ gdb <executable>
or
- \$ gdb
- (gdb) file <executable>

Using GDB

3. Run Program

- (gdb) run or
- (gdb) run [arguments]

4. In GDB Interactive Shell

- Tab to Autocomplete, up-down arrows to recall history
- help [command] to get more info about a command

5. Exit the gdb Debugger

- (gdb) quit

Run-Time Errors

- Segmentation fault
 - Program received signal SIGSEGV, Segmentation fault.
0x0000000000400524 in *function* (arr=0x7ffc902a270,
r1=2, c1=5, r2=4, c2=6) at *file.c*:12
 - Line number where it crashed and parameters to the function that caused the error
- Logic Error
 - Program will run and exit successfully
- How do we find bugs?

Setting Breakpoints

- Breakpoints
 - used to stop the running program at a specific point
 - If the program reaches that location when running, it will pause and prompt you for another command
- Example:
 - (gdb) break file1.c:6
 - Program will pause when it reaches line 6 of file1.c
 - (gdb) break my_function
 - Program will pause at the first line of my_function every time it is called
 - (gdb) break [position] if expression
 - Program will pause at specified position only when the expression evaluates to true

Breakpoints

- Setting a breakpoint and running the program will stop program where you tell it to
- You can set as many breakpoints as you want
 - (gdb) info breakpoints|break|br|b shows a list of all breakpoints

Deleting, Disabling and Ignoring BPs

- (gdb) delete [*bp_number* | *range*]
 - Deletes the specified breakpoint or range of breakpoints
- (gdb) disable [*bp_number* | *range*]
 - Temporarily deactivates a breakpoint or a range of breakpoints
- (gdb) enable [*bp_number* | *range*]
 - Restores disabled breakpoints
- If no arguments are provided to the above commands, all breakpoints are affected!!
- (gdb) ignore *bp_number iterations*
 - Instructs GDB to pass over a breakpoint without stopping a certain number of times.
 - *bp_number*: the number of a breakpoint
 - Iterations: the number of times you want it to be passed over

Displaying Data

- Why would we want to interrupt execution?
 - to see data of interest at run-time:
 - (gdb) `print` [/`format`] `expression`
 - Prints the value of the specified expression in the specified format
 - Formats:
 - d: Decimal notation (default format for integers)
 - x: Hexadecimal notation
 - o: Octal notation
 - t: Binary notation

Resuming Execution After a Break

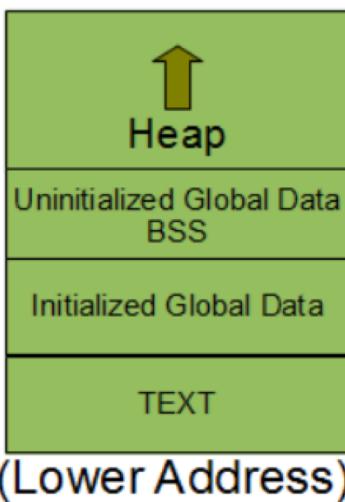
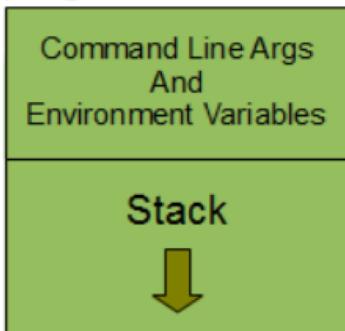
- When a program stops at a breakpoint
 - 4 possible kinds of gdb operations:
 - **c or continue**: debugger will continue executing until next breakpoint
 - **s or step**: debugger will continue to next source line
 - **n or next**: debugger will continue to next source line in the current (innermost) stack frame
 - **f or finish**: debugger will resume execution until the current function returns. Execution stops immediately after the program flow returns to the function's caller
 - the function's return value and the line containing the next statement are displayed

Watchpoints

- Watch/observe changes to variables
 - (gdb) `watch my_var`
 - sets a watchpoint on `my_var`
 - the debugger will stop the program when the value of `my_var` changes
 - old and new values will be printed
 - (gdb) `rwatch expression`
 - The debugger stops the program whenever the program reads the value of any object involved in the evaluation of `expression`

Process Memory Layout

(Higher Address)



(Lower Address)

- TEXT segment
 - Contains machine instructions to be executed
- Global Variables
 - Initialized
 - Uninitialized
- Heap segment
 - Dynamic memory allocation
 - malloc, free
- Stack segment
 - Push frame: Function invoked
 - Pop frame: Function returned
 - Stores
 - Local variables
 - Return address, registers, etc
- Command Line arguments and Environment Variables

Stack Info

- A program is made up of one or more functions which interact by calling each other
- Every time a function is called, an area of memory is set aside for it. This area of memory is called a **stack frame** and holds the following crucial info:
 - storage space for all the local variables
 - the memory address to return to when the called function returns
 - the arguments, or parameters, of the called function
- Each function call gets its own stack frame. Collectively, all the stack frames make up the **call stack**

Stack Frames and the Stack

```
1 #include <stdio.h>
2 void first_function(void);
3 void second_function(int);
4
5 int main(void)
6 {
7     printf("hello world\n");
8     first_function();
9     printf("goodbye goodbye\n");
10
11    return 0;
12 }
13
14 void first_function(void)
15 {
16     int imidate = 3;
17     char broiled = 'c';
18     void *where_prohibited = NULL;
19
20     second_function(imidate);
21     imidate = 10;
22 }
23
24
25 void second_function(int a)
26 {
27     int b = a;
28 }
```

Frame for `main()`

Frame for `first_function()`

Return to `main()`, line 9
Storage space for an int
Storage space for a char
Storage space for a `void *`

Frame for `second_function()`

Return to `first_function()`, line 22
Storage space for an int
Storage for the int parameter named `a`

When `first_function()` returns, its frame is used to determine where to return to (line 9 of `main()`) then, if `second_function()` has been deallocated and returned to the stack, the `void *` stored in `where_prohibited` is used to determine where to return to within `second_function()`.

Analyzing the Stack in GDB

- (gdb) backtrace | bt
 - Shows the call trace (the call stack)
 - Without function calls:
 - #0 main () at program.c:10
 - one frame on the stack, numbered 0, and it belongs to main()
 - After call to function display()
 - #0 display (z=5, zptr=0xbffffb34) at program.c:15
 - #1 0x08048455 in main () at program.c:10
 - Two stack frames: frame 1 belonging to main() and frame 0 belonging to display().
 - Each frame listing gives
 - the arguments to that function
 - the line number that's currently being executed within that frame

Analyzing the Stack

- **(gdb) info frame**
 - Displays information about the current stack frame, including its return address and saved register values
- **(gdb) info locals**
 - Lists the local variables of the function corresponding to the stack frame, with their current values
- **(gdb) info args**
 - List the argument values of the corresponding function call

Other Useful Commands

- (gdb) info functions
 - Lists all functions in the program
- (gdb) list
 - Lists source code lines around the current line

Lab 4

- Download old version of coreutils with buggy ls program
 - Untar, configure, make
- Bug: ls -t mishandles files whose time stamps are very far in the past. It seems to act as if they are in the future

```
$ tmp=$(mktemp -d)
$ cd $tmp
$ touch -d '1918-11-11 11:00 GMT' wwi-armistice
$ touch now
$ sleep 1
$ touch now1
$ TZ=UTC0 ls -lt --full-time wwi-armistice now now1
-rw-r--r-- 1 eggert csfac 0 1918-11-11 11:00:00.000000000 +0000 wwi-
armistice
-rw-r--r-- 1 eggert csfac 0 2017-01-25 00:11:55.528846902 +0000 now1
-rw-r--r-- 1 eggert csfac 0 2017-01-25 00:11:54.524820127 +0000 now
$ cd
$ rm -fr $tmp
```

Goal: Fix the Bug

- **Reproduce the Bug**
 - Follow steps on lab web page
- **Simplify input**
 - Run ls with –l and –t options only
- **Debug**
 - Use gdb to figure out what's wrong
 - \$ gdb ./ls
 - (gdb) run –lt /tmp/wwi-armistice /tmp/now /tmp/now1
(run from the directory where the compiled ls lives)
- **Patch**
 - Construct a patch “lab4.diff” containing your fix
 - It should contain a ChangeLog entry followed by the output of diff -u

Lab Hints

- Don't forget to answer all questions! (lab4.txt)
- Make sure not to submit a reverse patch! (lab4.diff)
- “Try to reproduce the problem in your home directory, instead of the \$tmp directory. How well does SEASnet do?”
 - Timestamps represented as seconds since Unix Epoch
 - SEASnet NFS filesystem has unsigned 32-bit time stamps
 - Local File System on Linux server has signed 32-bit time stamps
 - If you touch the files on the NFS filesystem it will return timestamp around 2054
 - => files have to be touched on local filesystem (df -l)
- Use “info functions” to look for relevant starting point